

Floating-Point Arithmetic

ECS130 Winter 2017

January 27, 2017

Floating point numbers

- ▶ Floating-point representation of numbers (scientific notation) has four components, for example,

$$\begin{array}{ccccccc} - & 3.1416 & \times 10^1 & \leftarrow \text{exponent} \\ \uparrow & \uparrow & \uparrow & \\ \text{sign} & \text{significand} & \text{base} & \end{array}$$

- ▶ Computers use binary (base 2) representation, each digit is the integer 0 or 1, e.g.

$$\begin{aligned} 10110 &= 1 \times 2^4 + 0 \times 2^3 + 1 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 \\ &= 22_{\text{in base 10}} \end{aligned}$$

and

$$\begin{aligned} \frac{1}{10} &= \frac{1}{16} + \frac{1}{32} + \frac{1}{256} + \frac{1}{512} + \frac{1}{4096} + \frac{1}{8192} + \dots \\ &= \frac{1}{2^4} + \frac{1}{2^5} + \frac{1}{2^8} + \frac{1}{2^9} + \frac{1}{2^{12}} + \frac{1}{2^{13}} + \dots \\ &= 1.100110011\dots \times 2^{-4} \end{aligned}$$

Floating point numbers

- ▶ The representation of a floating point number:

$$x = \pm b_0.b_1b_2 \cdots b_{p-1} \times 2^E$$

where

- ▶ It is *normalized*, i.e., $b_0 = 1$ (the hidden bit)
 - ▶ *Precision* ($= p$) is the number of bits in the **significand (mantissa)** (including the hidden bit).
 - ▶ *Machine epsilon* $\epsilon_m = 2^{-(p-1)}$, the gap between the number 1 and the smallest floating point number that is greater than 1.
- ▶ Special numbers
 $0, -0, \text{Inf} = \infty, -\text{Inf} = -\infty, \text{NaN} = \text{"Not a Number"}.$


IEEE standard

- ▶ All computers designed since 1985 use the *IEEE Standard for Binary Floating-Point Arithmetic* (ANSI/IEEE Std 754-1985), represent each number as a binary number and use binary arithmetic.
- ▶ Essentials of the IEEE standard:
 1. consistent representation of floating-point numbers by all machines adopting the standard;
 2. correctly rounded floating-point operations, using various rounding modes;
 3. consistent treatment of exceptional situation such as division by zero.

Citation of 1989 Turing Award to William Kahan: “... *his devotion to providing systems that can be safely and robustly used for numerical computations has impacted the lives of virtually anyone who will use a computer in the future*”

IEEE single precision format

- ▶ **Single** format takes 32 bits (=4 bytes) long:



- ▶ It represents the number $(-1)^s \cdot (1.f) \times 2^{E-127}$
- ▶ The leading 1 in the fraction need not be stored explicitly since it is always 1 (*hidden bit*)
- ▶ Precision $p = 24$ and machine epsilon $\epsilon_m = 2^{-23} \approx 1.2 \times 10^{-7}$
- ▶ The “ $E - 127$ ” in the exponent is to avoid the need for storage of a sign bit.

$$E_{\min} = (00000001)_2 = (1)_{10}, E_{\max} = (11111110)_2 = (254)_{10}.$$

- ▶ The range of positive normalized numbers:


$$N_{\min} = 1.00 \dots 0 \times 2^{E_{\min}-127} = 2^{-126} \approx 1.2 \times 10^{-38}$$

$$N_{\max} = 1.11 \dots 1 \times 2^{E_{\max}-127} \approx 2^{128} \approx 3.4 \times 10^{38}.$$

- ▶ Special representations for 0, $\pm\infty$ and NaN.

IEEE double precision format

- ▶ IEEE **double** format takes 64 bits (= 8 bytes) long:



- ▶ It represents the number $(-1)^s \cdot (1.f) \times 2^{E-1023}$
- ▶ Precision $p = 53$ and machine epsilon $\epsilon_m = 2^{-52} \approx 2.2 \times 10^{-16}$
- ▶ The range of positive normalized numbers is from

$$N_{\min} = 2^{-1022} \approx 2.2 \times 10^{-308}$$

$$N_{\max} = 1.11 \dots 1 \times 2^{1023} \approx 2^{1024} \approx 1.8 \times 10^{308}.$$

- ▶ Special representations for 0, $\pm\infty$ and NaN.

Rounding modes

- ▶ Let a positive real number x be in the normalized range, i.e., $N_{\min} \leq x \leq N_{\max}$, and write in the normalized form

$$x = 1.b_1b_2 \cdots b_{p-1}b_pb_{p+1} \cdots \times 2^E,$$

- ▶ Then the closest floating-point number less than or equal to x is

$$x_- = 1.b_1b_2 \cdots b_{p-1} \times 2^E$$

i.e., x_- is obtained by *truncating*.

- ▶ The next floating-point number bigger than x_- (also the next one that bigger than x) is

$$x_+ = (1.b_1b_2 \cdots b_{p-1} + 0.00 \cdots 01) \times 2^E$$

- ▶ If x is negative, the situation is reversed.

rounding modes, cont'd

Four rounding modes:

1. *round down*: $\text{fl}(x) = x_-$

2. *round up*: $\text{fl}(x) = x_+$

3. *round towards zero*:

$$\text{fl}(x) = x_- \text{ of } x \geq 0$$

$$\text{fl}(x) = x_+ \text{ of } x \leq 0$$

4. *round to nearest (IEEE default rounding mode)*:

$$\text{fl}(x) = x_- \quad \text{or} \quad x_+ \quad \text{whichever is nearer to } x.$$

Note: except that if $x > N_{\max}$, $\text{fl}(x) = \infty$, and if $x < -N_{\max}$, $\text{fl}(x) = -\infty$. In the case of tie, i.e., x_- and x_+ are the same distance from x , the one with its least significant bit equal to zero is chosen.

Rounding error

- ▶ When the *round to nearest* is in effect,

$$\text{relerr}(x) = \frac{|\text{fl}(x) - x|}{|x|} \leq \frac{1}{2}\epsilon_m.$$

- ▶ Therefore, we have

$$\text{relerr} = \begin{cases} \frac{1}{2} \cdot 2^{1-24} = 2^{-24} \approx 5.96 \cdot 10^{-8}, & \text{single precision} \\ \frac{1}{2} \cdot 2^{-52} \approx 1.11 \times 10^{-16}, & \text{double precision.} \end{cases}$$

Floating-point arithmetic

- ▶ IEEE rules for correctly rounded floating-point operations:

if x and y are correctly rounded floating-point numbers, then

$$\text{fl}(x + y) = (x + y)(1 + \delta)$$

$$\text{fl}(x - y) = (x - y)(1 + \delta)$$

$$\text{fl}(x \times y) = (x \times y)(1 + \delta)$$

$$\text{fl}(x/y) = (x/y)(1 + \delta)$$

where

$$|\delta| \leq \frac{1}{2}\epsilon_m$$

for the round to nearest,

- ▶ IEEE standard also requires that correctly rounded remainder and square root operations be provided.

Floating-point arithmetic, cont'd

IEEE standard response to exceptions

Event	Example	Set result to
Invalid operation	$0/0, 0 \times \infty$	NaN
Division by zero	Finite nonzero/0	$\pm\infty$
Overflow	$ x > N_{\max}$	$\pm\infty$ or $\pm N_{\max}$
underflow	$x \neq 0, x < N_{\min}$	$\pm 0, \pm N_{\min}$ or subnormal
Inexact	whenever $\text{fl}(x \circ y) \neq x \circ y$	correctly rounded value

Floating-point arithmetic error

- ▶ Let \hat{x} and \hat{y} be the floating-point numbers and that

$$\hat{x} = x(1 + \tau_1) \quad \text{and} \quad \hat{y} = y(1 + \tau_2), \quad \text{for } |\tau_i| \leq \tau \ll 1$$

where τ_i could be the relative errors in the process of “collecting/getting” the data from the original source or the previous operations, and

- ▶ **Question: how do the four basic arithmetic operations behave?**

Floating-point arithmetic error: $+$, $-$

Addition and subtraction:

$$\begin{aligned}\text{fl}(\hat{x} + \hat{y}) &= (\hat{x} + \hat{y})(1 + \delta), & |\delta| &\leq \frac{1}{2}\epsilon_m \\ &= x(1 + \tau_1)(1 + \delta) + y(1 + \tau_2)(1 + \delta) \\ &= x + y + x(\tau_1 + \delta + O(\tau\epsilon_m)) + y(\tau_2 + \delta + O(\tau\epsilon_m)) \\ &= (x + y) \left(1 + \frac{x}{x + y}(\tau_1 + \delta + O(\tau\epsilon_m)) + \frac{y}{x + y}(\tau_2 + \delta + O(\tau\epsilon_m)) \right) \\ &\equiv (x + y)(1 + \hat{\delta}),\end{aligned}$$

where $\hat{\delta}$ can be bounded as follows:

$$|\hat{\delta}| \leq \frac{|x| + |y|}{|x + y|} \left(\tau + \frac{1}{2}\epsilon_m + O(\tau\epsilon_m) \right).$$

Floating-point arithmetic error: $+$, $-$

Three possible cases:

1. If x and y have the same sign, i.e., $xy > 0$, then $|x + y| = |x| + |y|$; this implies

$$|\hat{\delta}| \leq \tau + \frac{1}{2}\epsilon_m + O(\tau\epsilon_m) \ll 1.$$

Thus $\text{fl}(\hat{x} + \hat{y})$ approximates $x + y$ well.

2. If $x \approx -y \Rightarrow |x + y| \approx 0$, then $(|x| + |y|)/|x + y| \gg 1$; this implies that $|\hat{\delta}|$ could be nearly or much bigger than 1.

This is so called **catastrophic cancellation**, it causes relative errors or uncertainties already presented in \hat{x} and \hat{y} to be magnified.

3. In general, if $(|x| + |y|)/|x + y|$ is not too big, $\text{fl}(\hat{x} + \hat{y})$ provides a good approximation to $x + y$.

Catastrophic cancellation: example 1

- ▶ Computing $\sqrt{x+1} - \sqrt{x}$ straightforward causes substantial loss of significant digits for large n

x	$\text{fl}(\sqrt{x+1})$	$\text{fl}(\sqrt{x})$	$\text{fl}(\text{fl}(\sqrt{x+1}) - \text{fl}(\sqrt{x}))$
1.00e+10	1.000000000049999994e+05	1.000000000000000000e+05	4.99999441672116518e-06
1.00e+11	3.16227766018419061e+05	3.16227766016837908e+05	1.58115290105342865e-06
1.00e+12	1.00000000000050000e+06	1.00000000000000000e+06	5.00003807246685028e-07
1.00e+13	3.16227766016853740e+06	3.16227766016837955e+06	1.57859176397323608e-07
1.00e+14	1.0000000000000503e+07	1.0000000000000000e+07	5.02914190292358398e-08
1.00e+15	3.16227766016838104e+07	3.16227766016837917e+07	1.86264514923095703e-08
1.00e+16	1.0000000000000000e+08	1.0000000000000000e+08	0.0000000000000000e+00

Catastrophic cancellation: example 1

- ▶ *Catastrophic cancellation can sometimes be avoided if a formula is properly reformulated.*
- ▶ For example, one can compute $\sqrt{x+1} - \sqrt{x}$ almost to full precision by using the equality

$$\sqrt{x+1} - \sqrt{x} = \frac{1}{\sqrt{x+1} + \sqrt{x}}.$$

Consequently, the computed results are

n	$\text{fl}(1/(\sqrt{n+1} + \sqrt{n}))$
1.00e+10	4.999999999875000e-06
1.00e+11	1.581138830080237e-06
1.00e+12	4.999999999998749e-07
1.00e+13	1.581138830084150e-07
1.00e+14	4.999999999999987e-08
1.00e+15	1.581138830084189e-08
1.00e+16	5.000000000000000e-09

Catastrophic cancellation: example 2

- ▶ Consider the function

$$h(x) = \frac{1 - \cos x}{x^2}$$

Note that $0 \leq f(x) < 1/2$ for all $x \neq 0$.

- ▶ Let $x = 1.2 \times 10^{-8}$, then the computed

$$\text{fl}(h(x)) = 0.770988\dots$$

is completely wrong!

- ▶ Alternatively, the function can be re-written as

$$h(x) = \left(\frac{\sin(x/2)}{x/2} \right)^2.$$

- ▶ Consequently, for $x = 1.2 \times 10^{-8}$, then the computed function $\text{fl}(h(x)) = 0.499999\dots < 1/2$ is fine!

Floating-point arithmetic error: $\times, /$

Multiplication and Division:

$$\begin{aligned}\text{fl}(\hat{x} \times \hat{y}) &= (\hat{x} \times \hat{y})(1 + \delta) \\ &= xy(1 + \tau_1)(1 + \tau_2)(1 + \delta) \\ &\equiv xy(1 + \hat{\delta}_\times), \\ \text{fl}(\hat{x}/\hat{y}) &= (\hat{x}/\hat{y})(1 + \delta) \\ &= (x/y)(1 + \tau_1)(1 + \tau_2)^{-1}(1 + \delta) \\ &\equiv xy(1 + \hat{\delta}_\div),\end{aligned}$$

where $\hat{\delta}_\times = \tau_1 + \tau_2 + \delta + O(\tau\epsilon_m)$

$\hat{\delta}_\div = \tau_1 - \tau_2 + \delta + O(\tau\epsilon_m)$.

Thus $|\hat{\delta}_\times| \leq 2\tau + \frac{1}{2}\epsilon_m + O(\tau\epsilon_m)$ and $|\hat{\delta}_\div| \leq 2\tau + \frac{1}{2}\epsilon_m + O(\tau\epsilon_m)$.

Multiplication and division are very well-behaved!

Reading

- ▶ Section 1.7 of *Numerical Computing with MATLAB* by C. Moler
- ▶ Websites discussions of numerical disasters:
 - ▶ T. Huckle, Collection of software bugs
<http://www5.in.tum.de/~huckle/bugse.html>
 - ▶ K. Vuik, Some disasters caused by numerical errors
<http://ta.twi.tudelft.nl/nw/users/vuik/wi211/disasters.html>
 - ▶ D. Arnold, Some disasters attributable to bad numerical computing
<http://www.ima.umn.edu/~arnold/disasters/disasters.html>
- ▶ In-depth material:
D. Goldberg, **What every computer scientist should know about floating-point arithmetic**, *ACM Computing Survey*, Vol.23(1), pp.5-48, 1991

LU factorization – revisited:

the need of pivoting in LU factorization, numerically

- ▶ LU factorization without pivoting

$$A = \begin{bmatrix} .0001 & 1 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 1 \end{bmatrix} = LU = \begin{bmatrix} 1 & \\ l_{21} & 1 \end{bmatrix} \begin{bmatrix} u_{11} & u_{12} \\ & u_{22} \end{bmatrix}$$

- ▶ In three decimal-digit floating-point arithmetic, we have

$$\begin{aligned}\hat{L} &= \begin{bmatrix} 1 & 0 \\ \text{fl}(1/10^{-4}) & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 10^4 & 1 \end{bmatrix}, \\ \hat{U} &= \begin{bmatrix} 10^{-4} & 1 \\ & \text{fl}(1 - 10^4 \cdot 1) \end{bmatrix} = \begin{bmatrix} 10^{-4} & 1 \\ & -10^4 \end{bmatrix},\end{aligned}$$

- ▶ Check:

$$\hat{L}\hat{U} = \begin{bmatrix} 1 & 0 \\ 10^4 & 1 \end{bmatrix} \begin{bmatrix} 10^{-4} & 1 \\ & -10^4 \end{bmatrix} = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 0 \end{bmatrix} \not\approx A,$$

LU factorization – revisited:

the need of pivoting in LU factorization, numerically

- ▶ Consider solving $Ax = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$ for x using this LU factorization.

- ▶ Solving

$$\hat{L}y = \begin{bmatrix} 1 \\ 2 \end{bmatrix} \implies \begin{array}{lcl} \hat{y}_1 & = & \text{fl}(1/1) = 1 \\ \hat{y}_2 & = & \text{fl}(2 - 10^4 \cdot 1) = -10^4. \end{array}$$

note that the value 2 has been “lost” by subtracting 10^4 from it.

- ▶ Solving

$$\hat{U}x = \hat{y} = \begin{bmatrix} 1 \\ -10^4 \end{bmatrix} \implies \begin{array}{lcl} \hat{x}_2 & = & \text{fl}((-10^4)/(-10^4)) = 1 \\ \hat{x}_1 & = & \text{fl}((1 - 1)/10^{-4}) = 0, \end{array}$$

- ▶ $\hat{x} = \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ is a **completely erroneous solution** to the correct answer $x \approx \begin{bmatrix} 1 \\ 1 \end{bmatrix}$.

LU factorization – revisited:

the need of pivoting in LU factorization, numerically

- ▶ LU factorization with partial pivoting

$$PA = LU,$$

- ▶ By exchanging the order of the rows of A , i.e.,

$$P = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix},$$

Then for the LU factorization of PA is given by

$$\begin{aligned}\hat{L} &= \begin{bmatrix} 1 & 0 \\ \text{fl}(10^{-4}/1) & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 10^{-4} & 1 \end{bmatrix}, \\ \hat{U} &= \begin{bmatrix} 1 & 1 \\ \text{fl}(1 - 10^{-4} \cdot 1) & 1 \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}.\end{aligned}$$

- ▶ The computed LU approximates A very accurately.
- ▶ As a result, the computed solution x is also perfect!